

Final Engagement

Attack, Defense & Analysis of a Vulnerable Network

Table of Contents

This document contains the following resources:



Network Topology & Critical Vulnerabilities



Alerts Implemented



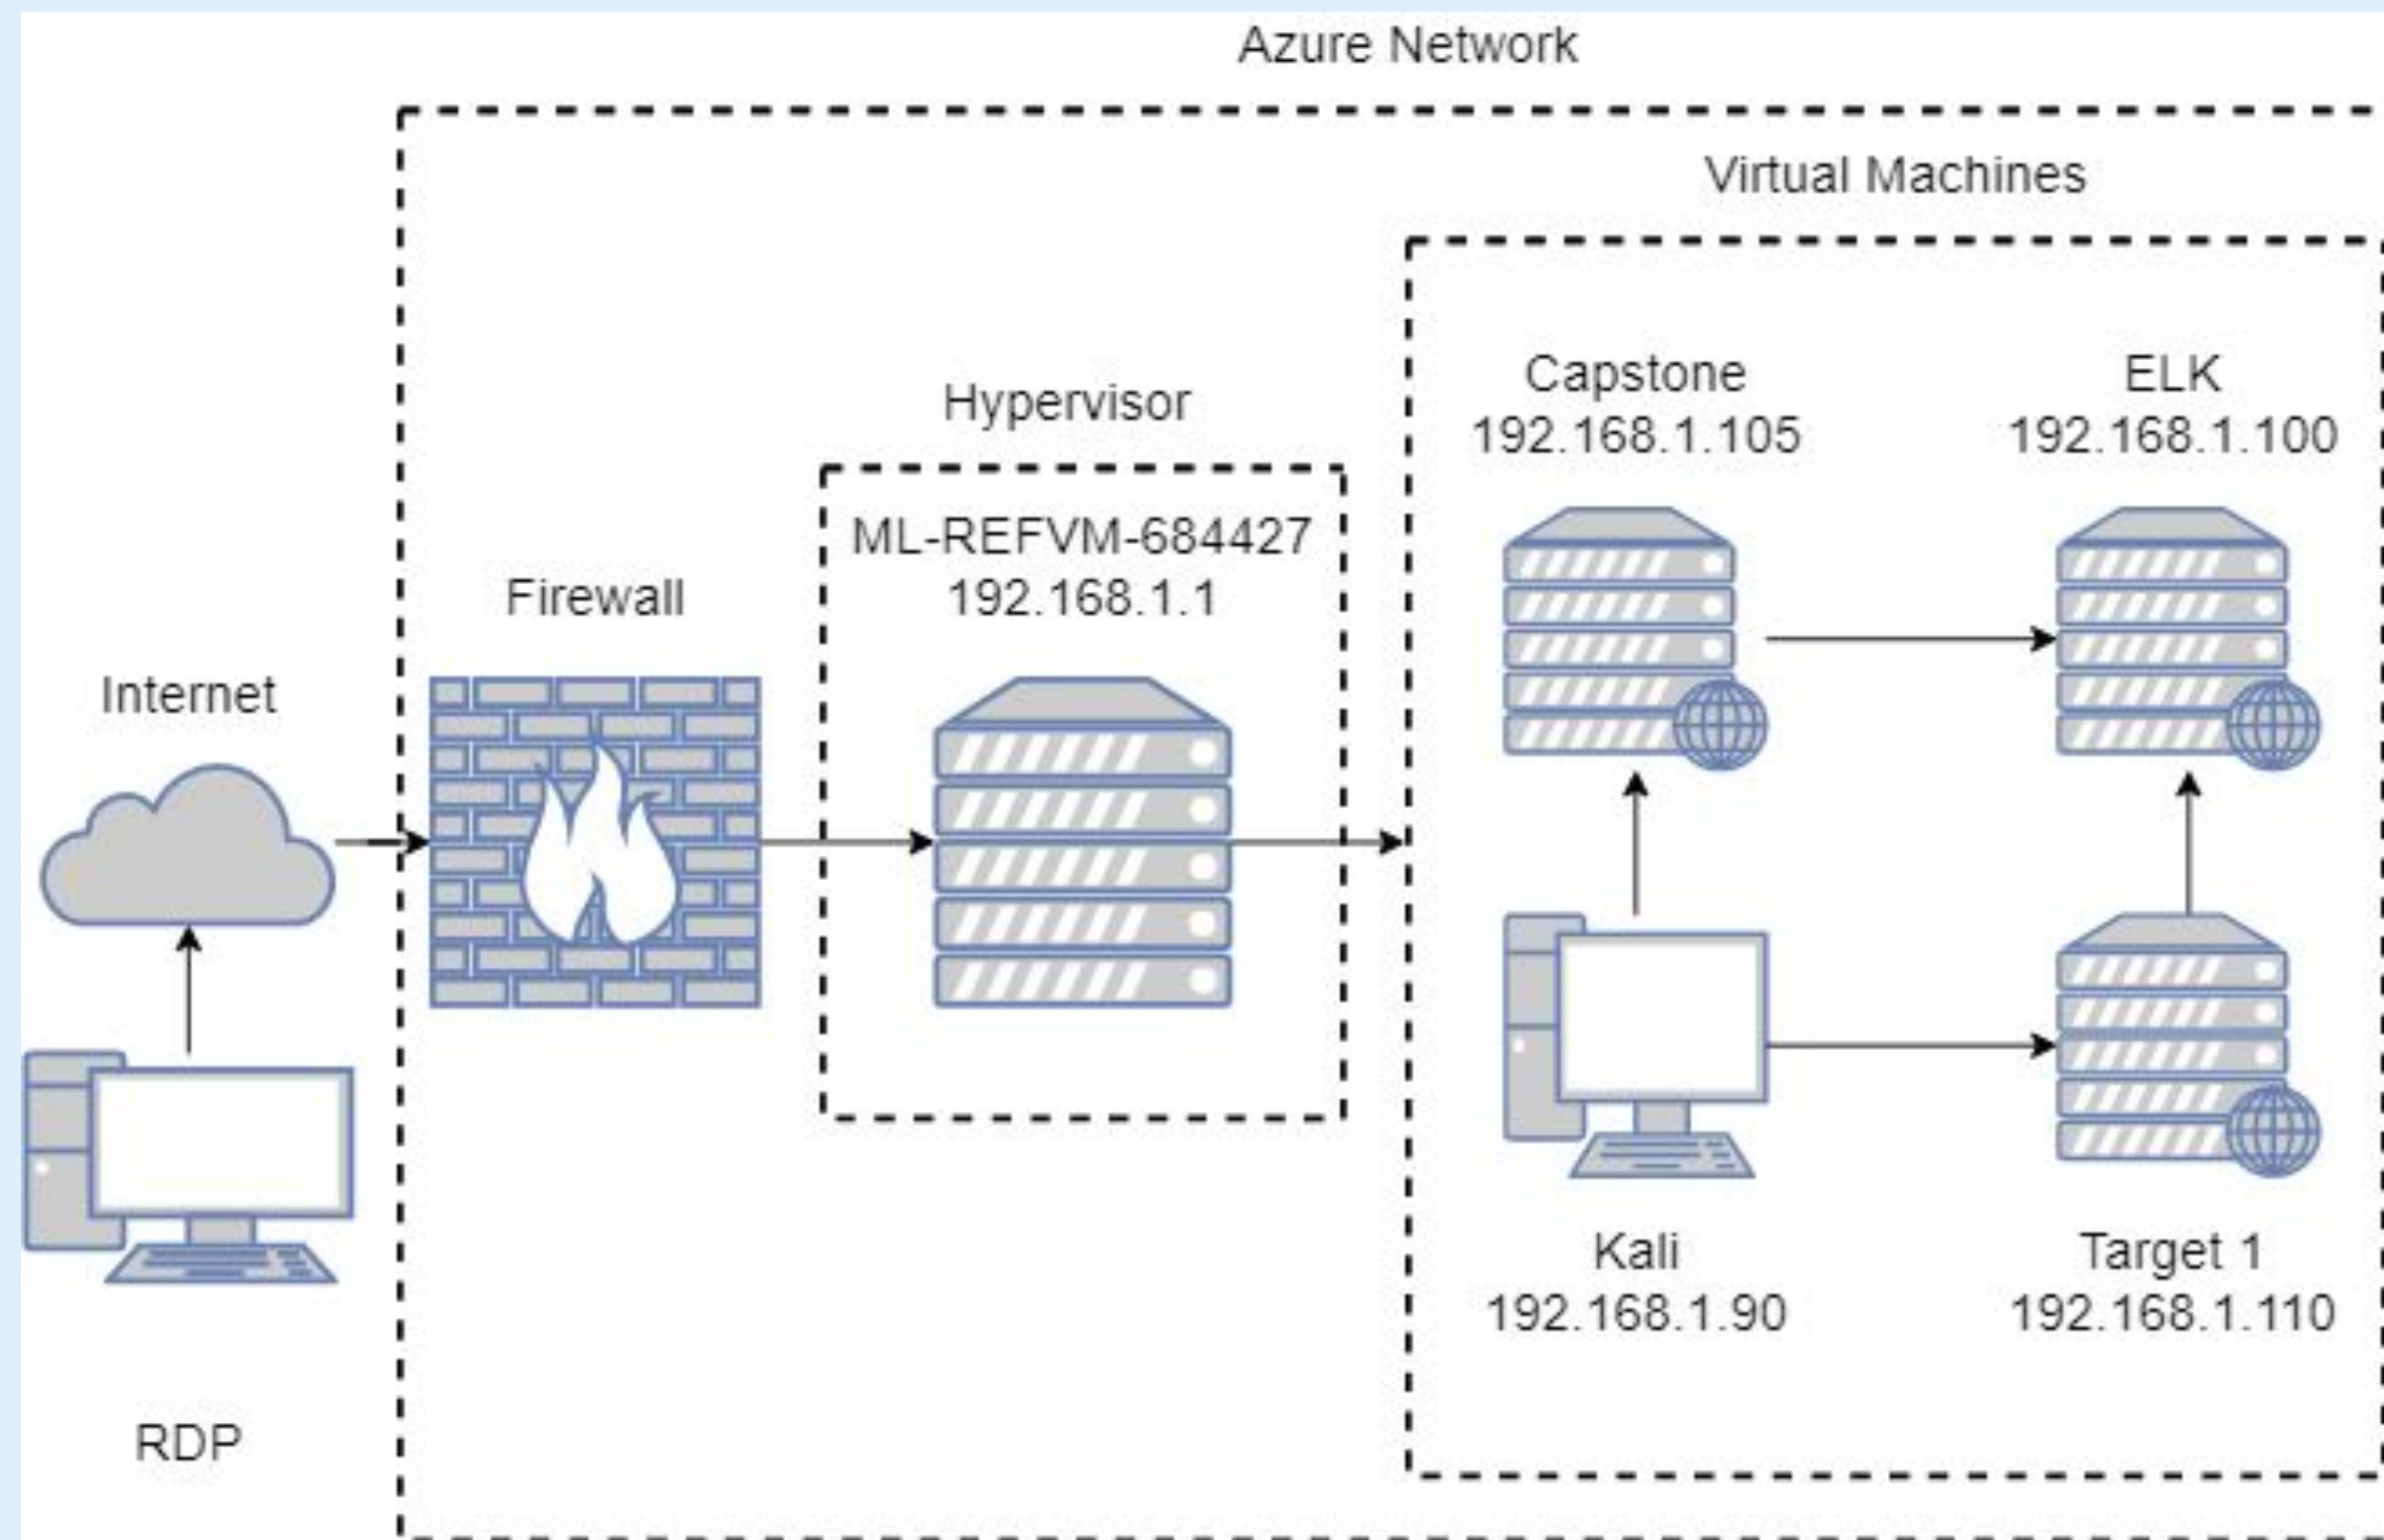
Hardening



Implementing Patches

Network Topology & Critical Vulnerabilities

Network Topology



Network

Address Range: 192.168.1.1/24

Netmask: 255.255.255.0

Gateway: 198.168.1.1

Machines

IPv4: 192.168.1.90

OS: Kali Linux

Hostname: Kali

IPv4: 198.168.1.110

OS: Linux

Hostname: Target 1

IPv4: 192.168.1.100

OS: Linux

Hostname: ELK

IPv4: 192.168.1.105

OS: Linux

Hostname: Capstone

Critical Vulnerabilities: Target 1

Our assessment uncovered the following critical vulnerabilities in **Target 1**.

Vulnerability	Description	Impact
Insecure Password Management	Redteam obtained passwords using dictionary brute force against the webserver and used credentials to SSH <ul style="list-style-type: none">- User Name: Micheal- Password: michael	This allows an attacker to gain access to protected web directories Note: Port 80 and & 22 were open
Privilege Escalation	We used Stevens sudo Python access to escalate from 'Steven to root'	Allowed privilege escalation to root. We used sudo -1 to gain information needed to perform escalation then sudo Python access to escalate to root
Wordpress User Enumeration	We utilized enum4linux to gather user information for the web server via SSH Users: michael/ steven/vagrant	Allows attacker to extract data and gather usernames to gain access to the web server
Unprotected and Unsalted Hash	Used JohnTheRipper to brute force the hash located within the MySQL database	Password hashing add a layer of security. Hashing allows passwords to be stored in a format that can't be reversed at any reasonable amount of time or cost for a hacker.



Alerts Implemented

Excessive HTTP Errors

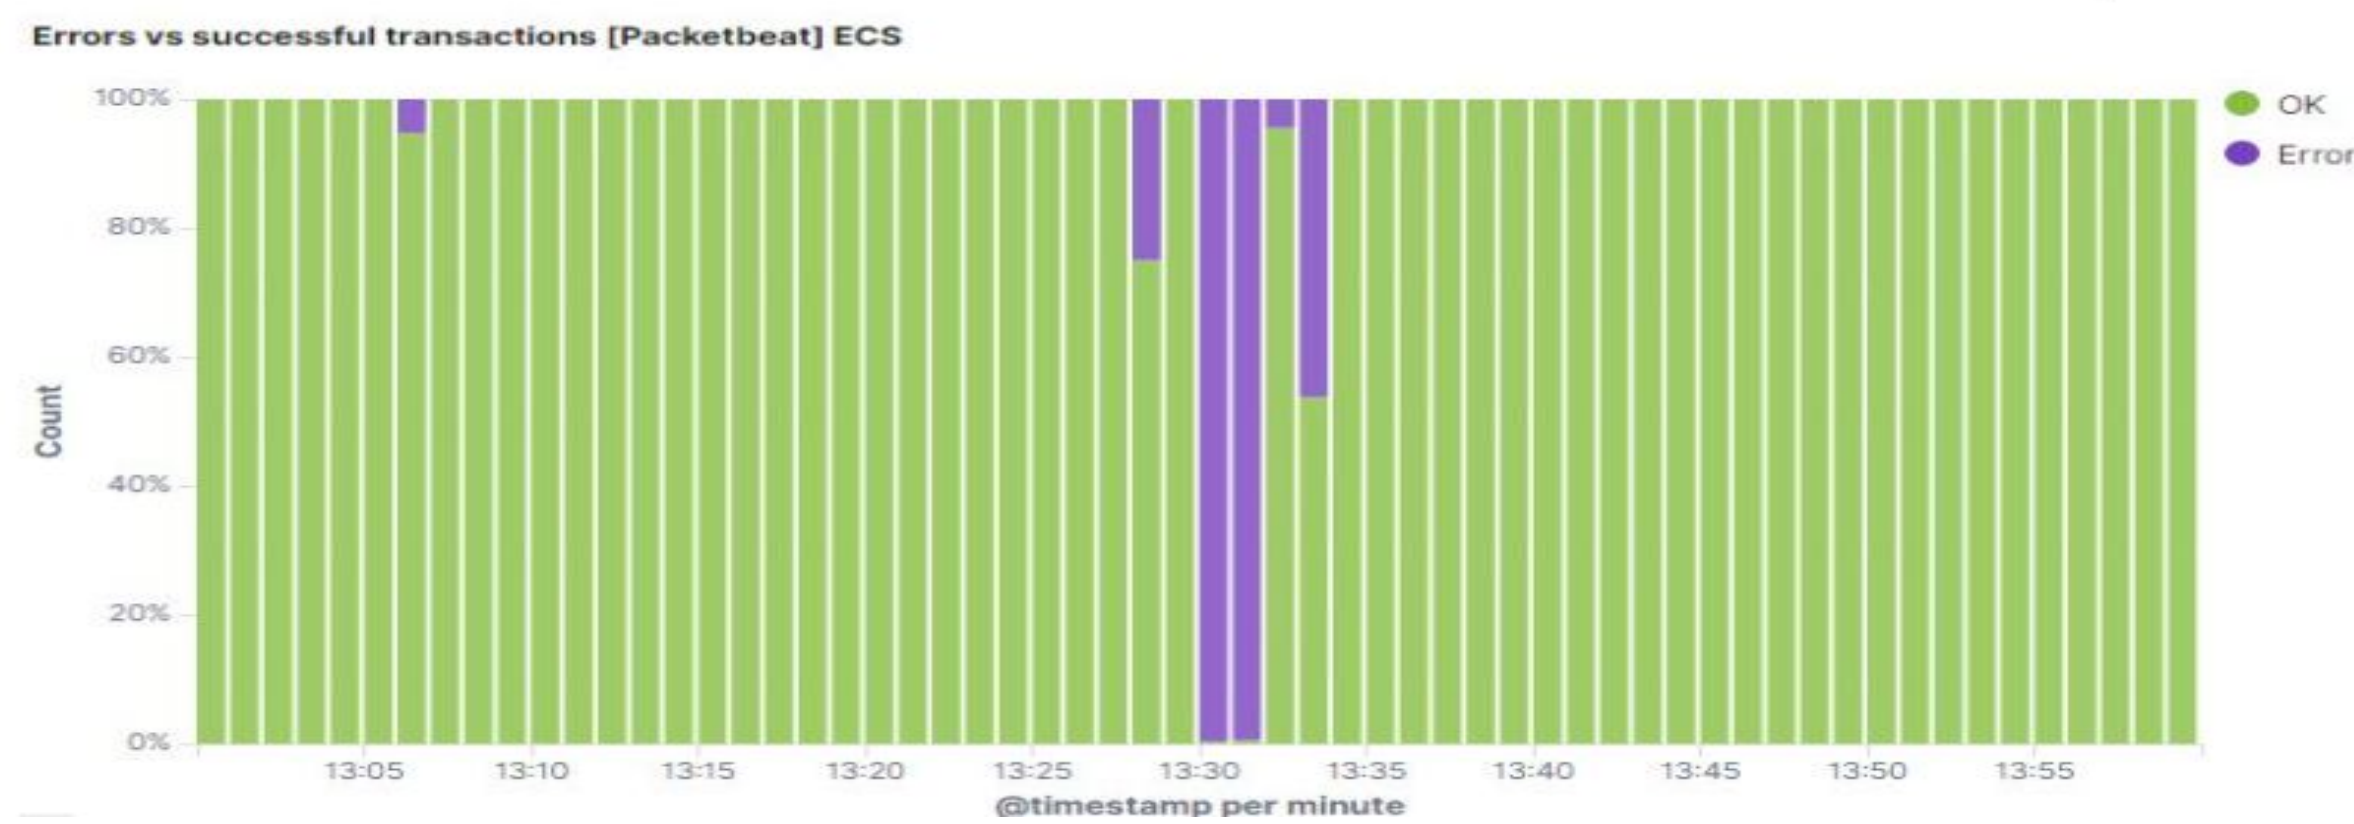
Summarize the following:

- Which **metric** does this alert monitor?
 - WHEN count() GROUPED OVER top 5 'http_response_status_code'
- What is the **threshold** it fires at?
 - Is ABOVE 400
- Provide a screenshot of the alert in action.

Excessive HTTP Errors

Summarize the following:

- This alert keeps track of any incoming HTTP error codes
- It is set to fire if more than 400 error codes are detected in the span of 5 minutes



Current status for 'Excessive HTTP Errors'

Execution history			Action statuses
Last one hour			
Trigger time	State	Comment	
2021-08-11T21:57:39+00:00	✓ OK		
2021-08-11T21:56:39+00:00	✓ OK		
2021-08-11T21:55:38+00:00	✓ OK		
2021-08-11T10:16:48+00:00	✓ OK		
2021-08-11T10:15:48+00:00	✓ OK		
2021-08-11T10:14:48+00:00	✓ OK		
2021-08-11T10:13:49+00:00	✓ OK		
2021-08-11T10:12:49+00:00	✓ OK		

HTTP Request Size Alert

- Metricbeat Indice: <WHEN sum() OF http.request.bytes OVER all documents IS ABOVE 3500 FOR THE LAST 1 minute>
- Metric alert: Monitoring HTTP Requests by size
- Threshold: Fires after any request that is over 3500 bytes in size.

HTTP Request Size Monitor

Summarize the following:

- This alert monitors the size of all incoming HTTP requests
- It is set up to fire upon receiving any HTTP request that is over 3500 bytes in size

Source IP	Destination IP	Source Bytes	Destination Bytes
192.168.1.90	192.168.1.110	5.1MB	18.5MB
192.168.1.90	192.168.1.115	89MB	177.9MB

Request Entity Too Large

The requested resource /wp-admin/update.php does not allow request data with POST requests, or the amount of data provided in the request exceeds the capacity limit.

WHEN sum() OF http.request.bytes OVER all documents IS ABOVE 3500 FOR THE LAST 1 minute


CPU Usage Monitor

metricbeat indice:

WHEN max() OF system.process.cpu.total.pct OVER all documents

IS ABOVE 0.5 FOR THE LAST 5 minutes

- **Metric:** Notify when the total of CPU processes
- **Threshold:** are at or are above 50% usage ----- in the last 5 minutes **(ALERT)**
- **Vulnerability Mitigated:** Malicious activity or malware in the midst of an attack. A spike in CPU usage can determine if someone is performing a *brute force attack*. This alert can work in combination of other metrics to determine illicit mining activity.
- **Reliability:** It's very reliable although it can create some potential false positives from multiple services running. However, that can help recognize and maintain background processes to improve CPU usage moving forward.




Current status for 'CPU Usage Monitor'

Execution history Action statuses

Last one hour

Trigger time	State
2021-08-11T23:17:49+00:00	✓ OK
2021-08-11T23:16:49+00:00	✓ OK
2021-08-11T23:15:49+00:00	✓ OK
2021-08-11T23:14:49+00:00	✓ OK
2021-08-11T23:13:49+00:00	✓ OK
2021-08-11T23:12:49+00:00	✓ OK
2021-08-11T23:11:49+00:00	✓ OK

Number of hosts [Met... CPU Usage Gauge [Me... Memory Usage Gauge ... Disk used [Metricbeat ... Inbound Traffic [Metri... Outbound Traffic [Met...



Top Hosts By CPU (Realtime) [Metricbeat System] ECS

server1		0.895%
target1		0.717%
target2		0.649%

Top Hosts By Memory (Realtime) [Metricbeat System] ECS

server1		9.422%
target1		8.862%
target2		8.707%

Hardening

Hardening Against Brute Force Attacks

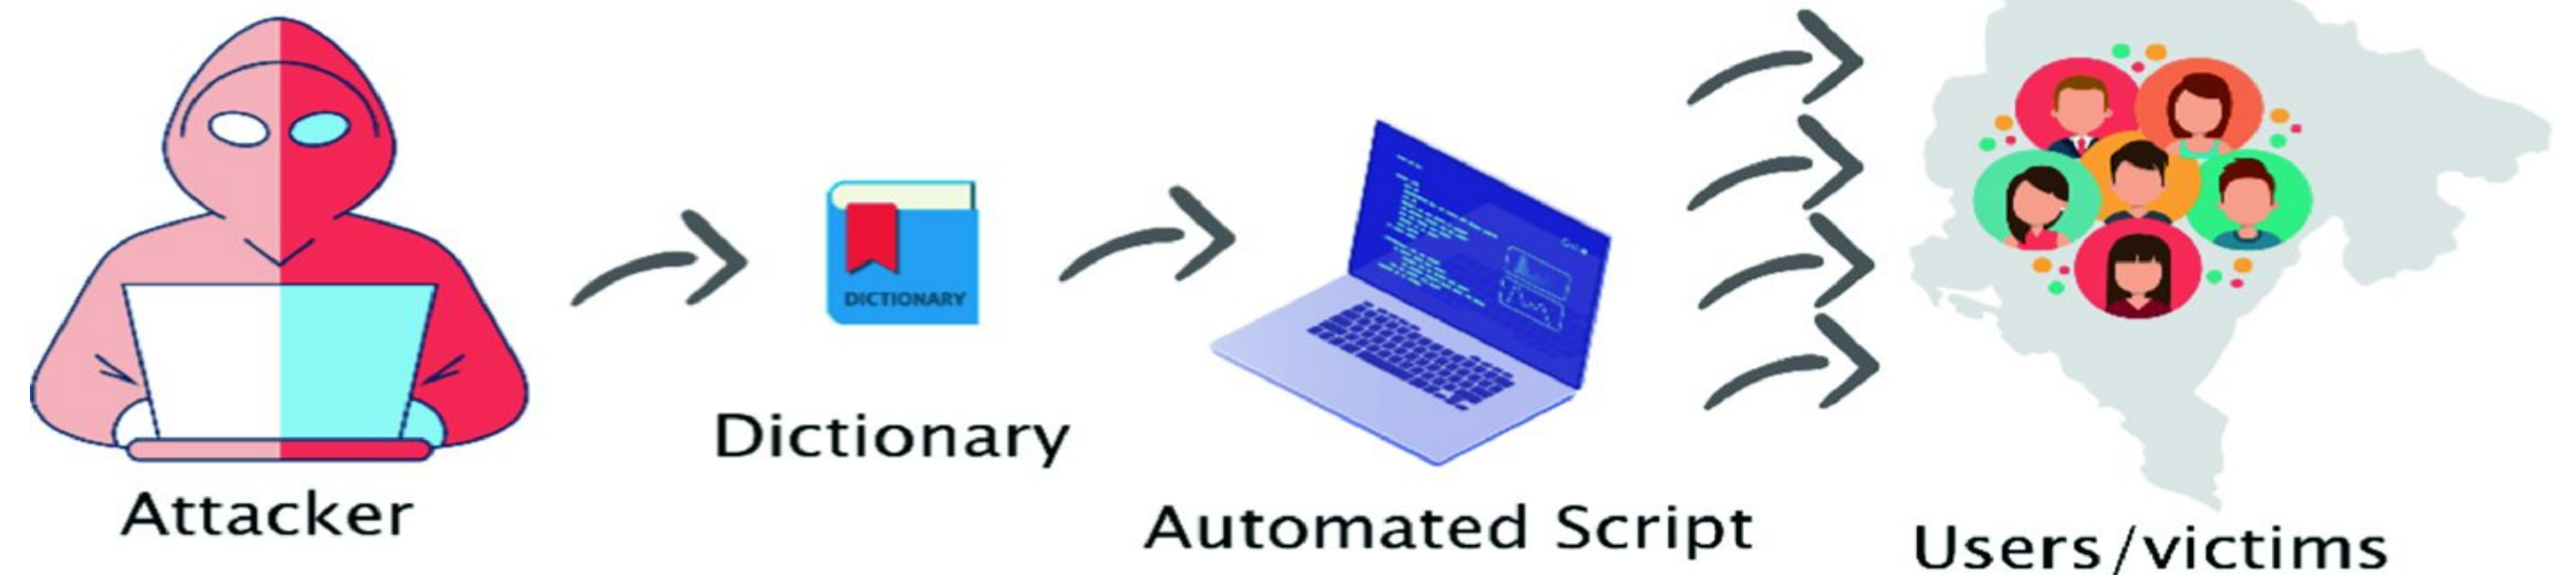
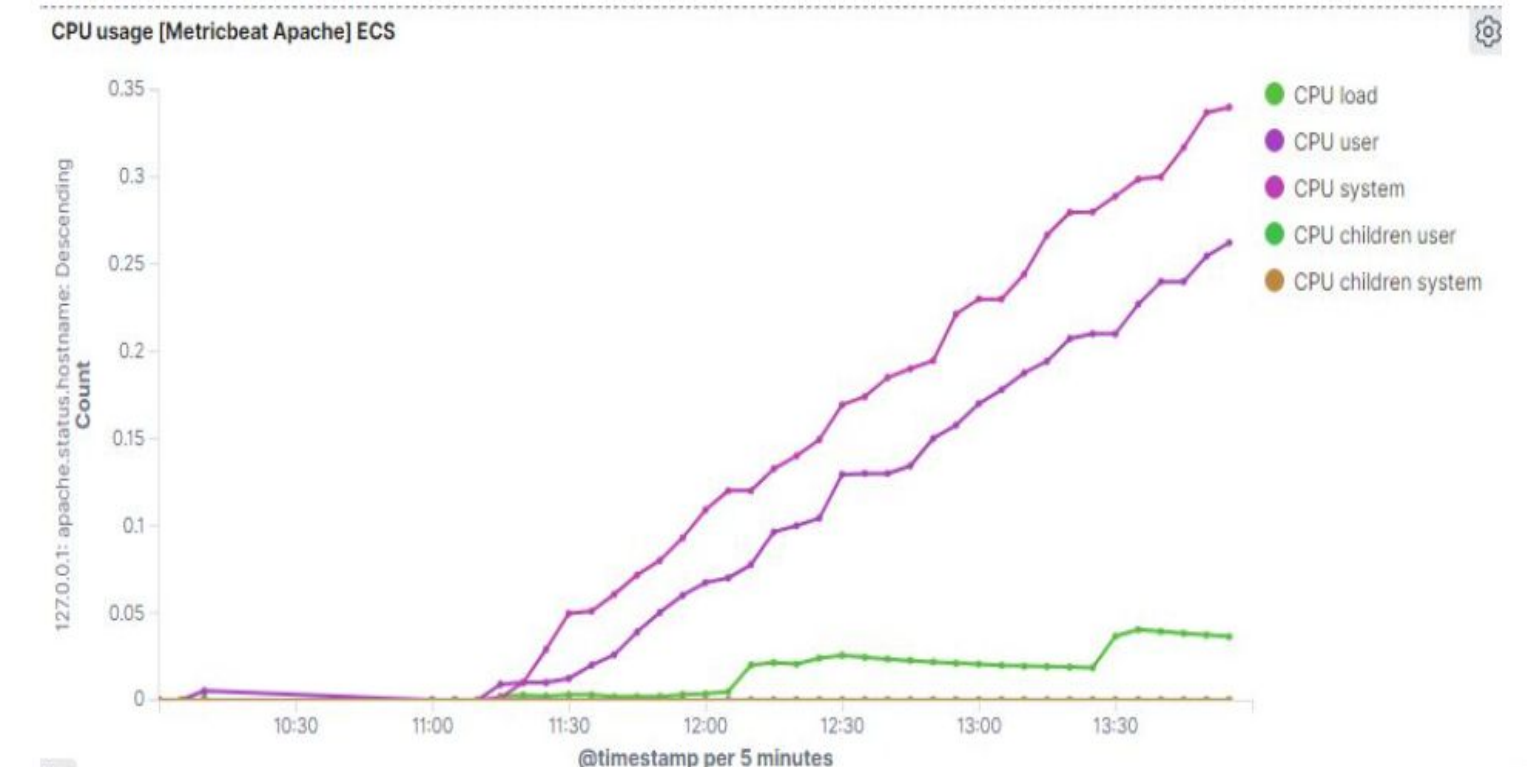
Creating a stronger password policy is a solid foundation of a hardening process. The policy will shield from brute force attacks. To implement this policy, we would access and edit the group policy, adding rules to the policy. We also will suggest the following actions:

- Account Lockout &/or Inject random pauses when checking a password to delay an attack
- Use CAPTCHA
- Ask users to answer a security response upon multiple failed login attempts.
- Having a strong, strict, password.
- Locking out IP addresses with multiple failed login attempts.
- Two factor authentication

CPU Usage Monitor

Summarize the following:

- Monitors percentage of total CPU usage since last event.
- Threshold fires if 50% is used within the last 5 minutes.



Hardening Against Port Scans

*Firewalls and IP tables are effective mitigation techniques against port scans.. Splunk and Kibana offer immediate alerting of the port scans which prompts a response from the security team. **We found that Port 80 and 22 were open!** To harden, the following steps are implemented.*

Mitigation of Port Scans

>**Block/forward/delay** port scans.

Create ipset lists:

```
<ipset create port_scanners hash:ip family inet hashsize 327 maxelem 65536 timeout 600>
```

```
<ipset create scanned_ports hash:ip, port family inet hashsize 327 maxelem 65536 timeout 60>
```

Create Rules:

```
< iptables -A INPUT -m state --state INVALID -j DROP
```

```
iptables -A INPUT -m state --state NEW -m set ! --match-set scanned_ports
```

```
src, dst -m hashlimit --hashlimit-above 1/hour --hashlimit-burst 5
```

```
--hashlimit-mode srcip --hashlimit-name portscan --hashlimithtable-expire
```

```
1000 -j SET --add-set port_scanners src --exist
```

```
iptables -A INPUT -m state --state NEW -m set --match-set port_scanners src -j DROP
```

```
iptables -A INPUT -m state --state NEW -j SET --add-set scanned_ports src, dst >
```

Firewall: Block all incoming and outgoing ports except for those needed

```
< iptables -A INPUT -p tcp -m multiport ! -dports 80, 443 -j DROP >
```

```
~ $ sudo nmap -v -sX -n 192.168.1.226 -p 22,80,445,5566,7878

Starting Nmap 7.01 ( https://nmap.org ) at 2019-01-29 19:33 EST
Initiating ARP Ping Scan at 19:33
Scanning 192.168.1.226 [1 port]
Completed ARP Ping Scan at 19:33, 0.23s elapsed (1 total hosts)
Initiating XMAS Scan at 19:33
Scanning 192.168.1.226 [5 ports]
Completed XMAS Scan at 19:33, 1.24s elapsed (5 total ports)
Nmap scan report for 192.168.1.226
Host is up (0.012s latency).
PORT      STATE       SERVICE
22/tcp    open|filtered ssh
80/tcp    closed      http
445/tcp    closed      microsoft-ds
5566/tcp   closed      westec-connect
7878/tcp   closed      unknown
MAC Address: 54:8C:A0:A4:00:0A (Unknown)

Read data files from: /usr/bin/../share/nmap
Nmap done: 1 IP address (1 host up) scanned in 1.61 seconds
Raw packets sent: 7 (268B) | Rcvd: 5 (188B)
```

> Iptables and ip6tables are used to set up, maintain, and inspect the tables of IPv4 and IPv6 packet filter rules in the Linux kernel. Several different tables may be defined. Each table contains a number of built-in chains and may also contain user-defined chains.

These rules assume any incoming connection that reaches the rule must be a result of a port scan, probably because all legitimate traffic has already been accepted. If there's more than 1 scanned missed port per hour from a source ip (the hashlimit matcher), the source ip is added to a cool down penalty box - the "port_scanners" ipset.

Hardening Against Wordpress User Enumeration



WordPress is a free, open-source website platform that is vulnerable to user enumeration by default.

An attacker could utilize a variety of methods

- **Error messages** could display if an existing user login is valid during a login attempt



- **REST API** provides data & content accessible to the public
- **Permalinks** are enabled. This allows the attacker to gain usernames by browsing through the author archives.

- `http://wordpress1.com/author/johnmcafee/`



- **Tools** to successfully scan, extract data & obtain user names.
 - **enum4linux - tool for enumerating information**
 - `enum4linux -a 192.168.1.110`
 - **wpscan - checks vulnerabilities & can enumerate users**
 - `wpscan -url http://192.168.1.110/wordpress - ue`

Wordpress User Enumeration Prevention



Web application firewall (WAF)

- This should be done by default.

Remove Detailed Error Messages that display valid usernames

Disable WordPress REST APIs (it's enabled by default) so that it can be inaccessible to the public.

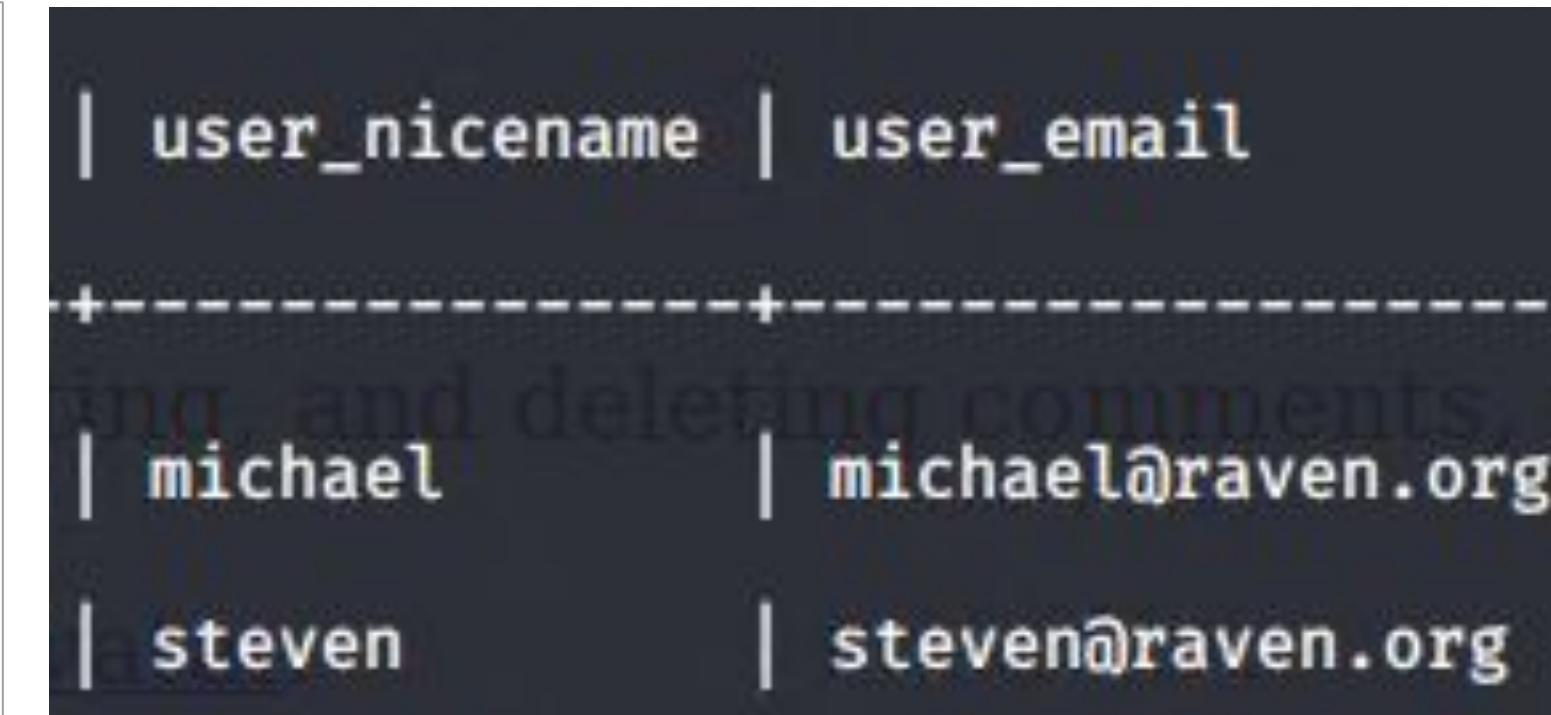
Author/user pages can be linked by iD rather than name or role (\$userid);

Mask `user_nicenames` by changing the value to something **random**

- This can be done by configuring the database. The random value could confuse the attacker once it's displayed to them
 - 'user_nicename' => 'charliebutterfly'

Additional security methods would be to use a WordPress security plugin as many are readily available.

- <https://wordpress.org/plugins/stop-user-enumeration/>
- <https://wordpress.org/plugins/disable-json-api/>
- <https://wordpress.org/plugins/edit-author-slug/>
- customize permalinks by changing /author/ portion of the URL
 - default: `http://example.com/author/username/`
 - custom: `http://example.com/cereal/chihuahua/`



Disable REST API

By default, this plugin ensures that the entire REST API is protected from non-authenticated users. You may use this page to specify which endpoints should be allowed to behave as normal.

IMPORTANT NOTE: Checking a box merely restores default functionality to an endpoint. Other authentication and/or permissions may still be required for access, or other themes/plugins may also affect access to those endpoints.

Rules for: Contributor

Manage Rules for Contributor Users

NOTE: New routes may be added in the future by plugins, themes, or WordPress itself. If you choose to manage access for a user role, you will have to come back and add permissions for any new routes later.

☒ Manage REST API Access ☐ Allow Full REST API Access

/ REST API ROOT

On this website, the REST API root is `http://dra2upgrade.local/wp-json/`

/oembed/1.0

☒ /oembed/1.0/embed

☐ /oembed/1.0/proxy

/wp/v2

☒ /wp/v2/posts

☒ /wp/v2/posts/(?P<id>[0-9]+)

☐ /wp/v2/posts/(?P<parent>[0-9]+)/revisions

☐ /wp/v2/posts/(?P<parent>[0-9]+)/revisions/(?P<id>[0-9]+)

☒ /wp/v2/posts/(?P<id>[0-9]+)/autosaves

☒ /wp/v2/posts/(?P<parent>[0-9]+)/autosaves/(?P<id>[0-9]+)


☐ /wp/v2/pages

☐ /wp/v2/pages/(?P<id>[0-9]+)

☐ /wp/v2/pages/(?P<parent>[0-9]+)/revisions

Implementing Patches

Implementing Patches with Ansible



Ansible for Security Infrastructure

Ansible is a configuration management tool that can be used to automate and patch systems. One can use Ansible to respond to threats and automate security solutions with modules, rules and playbooks.

Playbooks are written in YAML which allows for Human readable automation

Use Cases

- Infrastructure Provisioning
- Deploying security updates
- Installing updates on a test environment
- Identifying unknown issues
- Creating playbooks for configuration changes

```
- name: Update root password
  hosts: all
  become: yes

  vars:
    root_password: "{{ root_password }}"
    encrypt: "sha512_crypt"
    confirm: yes
    salt_size: 7

  tasks:
    - name: Update root password
      user:
        name: root
        password: "{{root_password | password_hash('sha512') }}"
        update_password: always
```

Can be scheduled to update root password every 60 days

Stay up to date and automate!

Resources:

OWASP - https://owasp.org/www-community/controls/Blocking_Brute_Force_Attacks

Common Vulnerabilities and Exposure (CVE) - https://cve.mitre.org/cve/search_cve_list.html

Palo Alto Networks - <https://docs.paloaltonetworks.com/>

Ansible - <https://www.ansible.com/use-cases/security-automation>

WordPress.org - <https://wordpress.org/plugins/>

Trustwave - <https://www.trustwave.com/en-us/resources/security-resources/security-advisories/>

Linux IP Tables - <https://www.man7.org/linux/man-pages/man8/iptables.8.html>



The End